

Contents

Contents	1
List of Figures	3
List of Tables	5
1 NEFI	7
1.1 abstract	7
1.2 introduction	7
1.3 Network Extraction From Images	10
1.3.1 Preprocessing Collection	11
1.3.2 Segmentation Collection	16
1.3.3 Graph Detection Collection	17
1.3.4 Graph Filter Collection	18
1.3.5 Defining a Graph Similarity Measure	19
Bibliography	21
Todo list	23

--

--

--

List of Figures

1.1	Pipeline components of NEFI	10
1.2	NEFI's pipeline executed	12
1.3	NEFI's graphical user interface.	13
1.4	<i>P. polycephalum</i> : Input and output of NEFI	14
1.5	<i>A. junius</i> : Output of NEFI	15

--

--

--

List of Tables

--

--

--

1 NEFI

1.1 abstract

Networks are amongst the central building blocks of many systems. Given a graph of a network, methods from graph theory enable a precise investigation of its properties. Software for the analysis of graphs is widely available and has been applied to study various types of networks. In some applications, graph acquisition is relatively simple. However, for many networks data collection relies on images where graph extraction requires domain-specific solutions. Here we introduce NEFI, a tool that extracts graphs from images of networks originating in various domains. Regarding previous work on graph extraction, theoretical results are fully accessible only to an expert audience and ready-to-use implementations for non-experts are rarely available or insufficiently documented. NEFI provides a novel platform allowing practitioners to easily extract graphs from images by combining basic tools from image processing, computer vision and graph theory. Thus, NEFI constitutes an alternative to tedious manual graph extraction and special purpose tools. We anticipate NEFI to enable time-efficient collection of large datasets. The analysis of these novel datasets may open up the possibility to gain new insights into the structure and function of various networks. NEFI is open source and available at <http://nefi.mpi-inf.mpg.de>.

1.2 introduction

The study of complex network-like objects is of increasing importance for many scientific domains. The mathematical study of networks, Graph Theory, formalizes a network's structure by modeling the constituents of a network as *vertices* and the pairwise relations between them as *edges*. Some communities traditionally refer to vertices as nodes or sites and to edges as arcs or links. Networks are ubiquitous in everyday life. Examples are as diverse as the Internet, social networks, transportation networks, metabolic networks, blood vessels or the vein networks of leaves. For a comprehensive review see [17].

In situations where the extraction of a mathematical graph from a physical network is easy, the size of graphs that can be analyzed quickly increased from hundreds to millions of vertices. At the same time it became feasible to build large

databases of various types of networks. This enabled the application of software incorporating methods from statistics and graph theory to obtain many results that changed our understanding of large scale network structures. However, digitization remains difficult for many types of networks, e.g. leaf venations, blood vessels or food webs, and therefore ready-to-analyze datasets are often not available. In these cases, investigation on a larger scale requires tedious and sometimes error prone data acquisition.

In many experimental settings networks are initially available as high quality images obtained under laboratory control. Before any analysis can take place, it is necessary to extract the associated graphs from these images. This requires the identification of vertices and edges within the depicted structure. This process can quickly become very work-intensive even for smaller networks, which makes automated solutions indispensable.

Leveraging advances in computer vision, several authors have proposed and successfully implemented solutions for domain specific graph extraction applications. The authors of [18], [19] consider the mycelial networks of *P. impudicus*. They use watershed segmentation in combination with a novel enhancement step designed to highlight curvilinear features in the input networks. Based on the segmented image a skeleton is computed and used to extract the graph representing the input network. The resulting method is designed to be brightness and contrast invariant in order to correctly extract the networks grown by *P. impudicus* from challenging noisy or low contrast images.

Baumgarten et al. [4], [5] investigate the vein networks of *P. polycephalum*. For segmenting the input image they rely on careful constant thresholding followed by a sequence of restoration algorithms that try to repair the network in the segmented image. Next, the restored segmented image is used to compute a skeleton. After applying another sequence of correction steps, the skeleton is scanned to extract the graph of the input network.

In [7] a more general algorithm applicable to a variety of problems is proposed. Based on an original stochastic model, the authors use Monte Carlo sampling to obtain junction-points in the input image. This technically involved solution guarantees structural coherence for the resulting graph representation. Further examples include the extraction of road networks [16], retinal blood vessel analysis [12] and the extraction of plane graphs [21].

The three above mentioned algorithmic solutions for the network extraction problem exhibit one or more of the following limitations:

- They do not build on top of well-established computer vision methods and tend to rely on ad-hoc algorithms. As a result the quality of the method and its implementation could likely be improved. In addition, a lot of time is spent on reimplementing algorithms that are already available.

- They are not implemented or only available as pseudo-code.
- They are implemented but not designed for easy of use, distribution and extendability.

We are aware that the primary objective of the work cited above is not the production of reusable software, but of algorithms and tools for solving a concrete research question. As a result, the above authors have limited time for researching advances in computer vision, following best software engineering practice or writing documentation respectively.

From experience we know that when producing an easy-to-use software, a large part of the required work consists of specifying and improving the user-interface as well as working out minor bugs and annoyances. This type of work, while time consuming, is essential for any software aiming to reach a non-negligible audience. However, efforts like these are hardly attractive to researchers whose focus is on obtaining the next result. While we understand that under these circumstances the aforementioned limitations arise naturally, we strongly believe that it is necessary to overcome those limitations in order to increase the value and the impact of scientific software in general and network extraction software in particular.

To this end, we introduce NEFI, a lightweight piece of ready-to-go software intended to enable the non-expert to automatically extract networks from images. NEFI constitutes an extensible framework of interchangeable algorithms accessible through an intuitive graphical user interface.

We emphasize at this point that we do not claim to introduce novel techniques for image processing or computer vision. Instead, our contribution consists of a reusable, flexible and easily extendable toolbox combining well-known methods, which have become standard in their respective fields of origin, in a meaningful way. By introducing NEFI, we hope to make these methods more widely accessible to practitioners in other fields.

NEFI's segmentation is based on a combination of standard routines available in [OpenCV](#), (2015), [6]. These algorithms are known to perform well on clean and uncluttered images obtained under controlled laboratory conditions. However, on more challenging inputs of low contrast, strong gradients or similar irregularities, their performance is severely reduced. Nevertheless, in these cases more involved algorithms, currently not implemented as part of a reliable library and thus not integrated into NEFI, may still be able to process these images. To help meet this situation, NEFI was designed with extendability in mind. As a result users will find it easy to build on-top of NEFI's code in order to add their own implementations of more sophisticated methods.




Figure 1.1: A flow chart illustrating NEFI's pipeline components in green boxes. Dashed arrows depict optional sections of the pipeline. Blue and orange boxes denote NEFI's input and possible outputs respectively.

1.3 Network Extraction From Images

NEFI features a collection of image processing routines, segmentation methods and graph algorithms designed to process 2D digital images of various networks and network-like structures. Its main function is executing a so-called extraction pipeline, designed to analyze the structures depicted in the input image. An extraction pipeline, for short pipeline, denotes an ordered sequence of algorithms. A successful execution will return a representation of the network in terms of an edge-weighted undirected planar graph. Computed weights include edge lengths and edge widths. Once the graph is obtained, available graph analysis software [2], [3], [10], [13], [14], [22] or custom written scripts can be deployed to investigate its properties.

A typical pipeline combines algorithms from up to four different classes: preprocessing, segmentation, graph detection and graph filtering, see Figure 1.1. A more detailed description follows below.

For each pipeline section, NEFI typically offers several interchangeable algorithms to choose from. After executing preprocessing routines, a segmentation algorithm separates foreground from background. Then the foreground is thinned

to a skeleton from which the vertices and edges of the graph are determined. In the process various edge weights are computed. Finally, the graph can be subjected to a variety of useful graph filters. Figure 1.2 illustrates the intermediate results of NEFI’s pipeline steps listed in the order of their execution. When a pipeline is executed, NEFI makes all intermediate results available via its clean and intuitive GUI, see Figure 1.3.

Using the GUI all basic functions of NEFI can be accessed in an intuitive fashion. To facilitate ease-of-use, most of NEFI’s algorithms come with default parameters based on the settings in [OpenCV](#) [6], which were found to perform well on our test sets as well as on many other images.

There are various predefined pipelines to get started immediately. Alternatively, users may freely combine the various methods to build custom pipelines. Both approaches allow the user to experiment with the available methods in order to close in on the optimal settings for the data. Once a pipeline is constructed, it can be saved and reused. NEFI’s simple pipeline concept together with a self-explanatory graphical user interface make working with NEFI intuitive and straightforward. NEFI also offers a command-line mode, which is suited for batch processing.

NEFI comes with a number of example images from different domains which we use to produce the figures in this work. Figure 1.4 and Figure 1.5 show NEFI’s output on two images using predefined pipelines. Blue squares denote the vertices and red lines the edges of the detected graph. The thickness of the detected edges corresponds to thickness of the depicted structures. For comparison the extracted graph is drawn on top of the input image. We present a detailed quantitative evaluation in a later section.

We stress that NEFI can deal with a range of inputs from various domains as long as they are of sufficient quality. In addition to the examples shown above, it has been successfully used to process images of natural (e.g. leaf venation, patterns of mud cracks) as well as man-made structures (tilings). It is also straightforward to add custom extensions. We provide a well documented platform which allows programmers to include more specialized segmentation algorithms or additional graph filters. For an overview of alternative graph extraction approaches see for example [8].

Next, we discuss the purpose and design of each major stage of the pipeline and highlight some of NEFI’s strong points.

1.3.1 Preprocessing Collection

The preprocessing section of the pipeline offers various standard image processing algorithms intended to be used prior to the segmentation step. Preprocessing methods may be exploited to affect the output of the segmentation step. For example, adding a slight blur to an input image may benefit the overall result

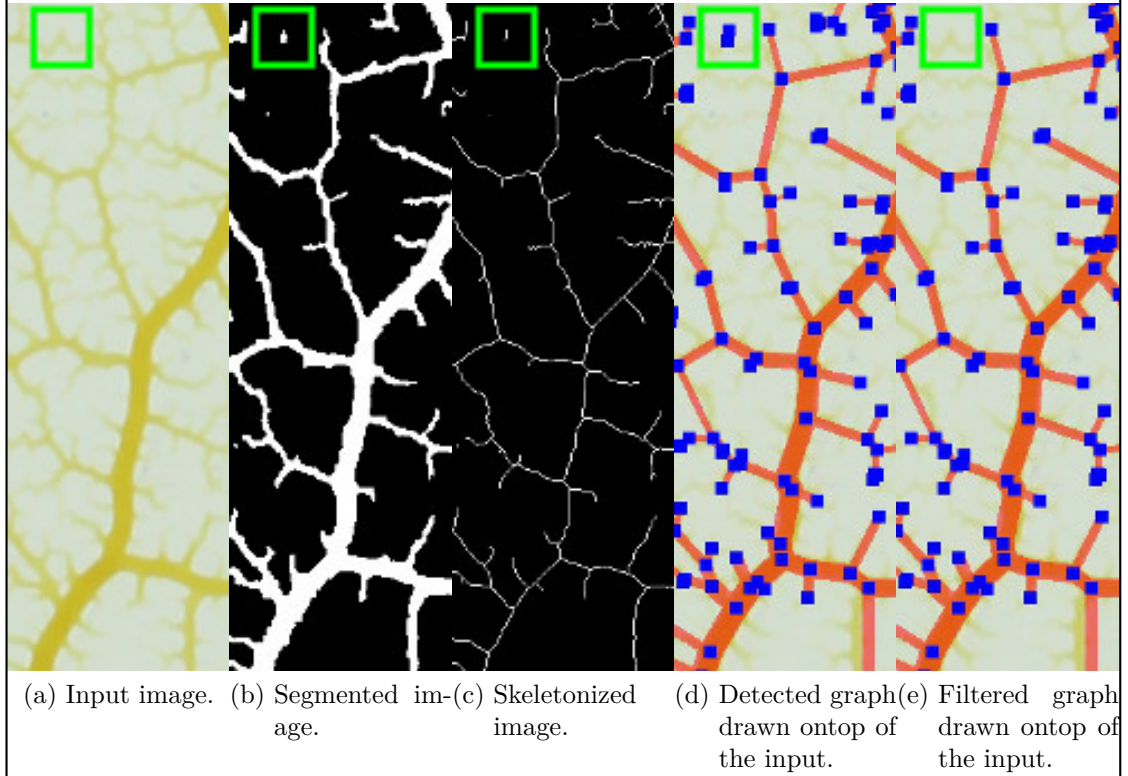


Figure 1.2: Direct comparison of NEFI's pipeline steps given a slice of an image of a slime mold (*P. polycephalum*). From left to right: input image, segmented image, skeletonized image, detected graph and filtered graph. The green square contains a very faint vein which the segmentation did not pick up fully. Thus, the skeleton becomes fragmented which leads to spurious vertices in the detected graph. By applying a graph filter we remove stray vertices without manipulation of the segmented or the skeletonized image. Similar filtering can remove “dead-ends”, i.e. vertices that do not belong to any cycle in the graph.




Figure 1.3: A screenshot of NEFI's GUI running on Mac OS. On the left hand side NEFI lists intermediate results as thumbnails. Bringing the final result to the center workspace allows for direct visual assessment of the quality of the extracted graph. On the right hand side NEFI's pipeline elements can be accessed via drop-down menus. The image was produced in a collaboration with the KIST Europe.




Figure 1.4: Extracted graph of the network formed by a slime mold (*P. polycephalum*). The left hand side shows the input image depicting the network. The right hand side shows the extracted graph overlayed on top off the same image for direct comparison. Note, that no filters have been applied. The image was produced in a collaboration with the KIST Europe.




Figure 1.5: Extracted graph of the vein network exhibited by a wing of a dragonfly (*A. junius*). Note, that after the use of various filters a very clean-looking graph is obtained. Image courtesy of Pam and Richard Winegar.

by reducing the amount of spurious white pixels in the segmented image. However, blurring too much will remove fine detail and reduce accuracy in determining the thickness of depicted edges. As a result, we recommend to experiment with different approaches and parameter settings in order to decide how to use preprocessing. For images of sufficient quality we found that excellent results can be obtained without preprocessing.

NEFI relies on [OpenCV](#) [6] for preprocessing and offers Gaussian and Median Blurring, Denoising as well as Bilateral Filtering.

1.3.2 Segmentation Collection

The goal of the segmentation step is separating the image foreground, i.e. the structures of interest, from the remaining image. NEFI builds on top of [OpenCV](#) [6] combining different segmentation algorithms. The general-purpose algorithms shipped with NEFI have become standard in image processing and perform reliably well if input images are clean, devoid of strong gradients and have a good contrast between fore- and background. Conversely, if the input becomes more challenging, the effectiveness of NEFI's segmentation degrades quickly and more complex or domain-specific algorithms become necessary. We defer a quantitative study of how the properties of the input image affect NEFI's performance to the Section *Evaluation*.

NEFI's segmentation is designed such that several algorithms can be used interchangeably. We included basic thresholding algorithms like Otsu's method [1] or adaptive thresholding as well as more involved segmentation routines such as guided watershed [15] and the GrabCut algorithm [20]. The last two methods receive as an additional input a so-called marker. The better the markers approximate the foreground, the better these algorithms work. NEFI offers several marker strategies which can be used interchangeably together with the respective marker based segmentation routines.

Interchangeability of the algorithms is a core design principle of all pipeline steps.

This design facilitates easy experimentation with different methods. Our own experience shows that often it is not clear a priori which methods work for a given input image. This decision usually also depends on the desired degree of detail in the final output, where less sensitive methods might produce fewer false positives. This ease of experimentation with quick visual feedback from the GUI is one of the major strong points of NEFI.

The flexibility is not limited to the algorithms we provide. Instead, NEFI's software design makes it easy to integrate additional methods. We expect that in practice challenging inputs will be encountered for which the algorithms currently offered by NEFI will be insufficient. In these cases a potential user may choose

to implement additional, perhaps domain-specific, methods. By extending NEFI, the user can rely on existing modules and thus save a lot of time. The authors are convinced that improved extendability is another strong point of our work.

1.3.3 Graph Detection Collection

The graph detection collection consists of algorithms that take a segmented image as input and detect the nodes and the edges of the graph. We offer a colloquial description of the actual algorithm because we do not rely on well-documented library code for this section of the pipeline.

The first step for graph detection is called thinning. Here we reduce the segmented foreground such that every line is only one pixel thick, while preserving the connectivity properties of both the foreground and the background pixels. The result of this process is called the skeleton of the segmented image. To do so we implemented the algorithm by Guo and Hall [9]. It always produces thin results and preserves 8-connectivity of the foreground pixels. A pure Python implementation proved to be fairly slow, hence we chose to implement this function as a C extension.

For fairly thin foreground features this method is nearly flawless and finds a skeleton where the lines lie in the center of the foreground areas. However, large foreground sections lead to artifacts in the skeleton whose exact shape depends on the noise present at their borders.

On the skeleton we then detect the positions of nodes. For this purpose we adapt criteria from thinning algorithm by Zhang and Suen [23]. A white pixel becomes a node if its removal creates exactly one or at least three 4-connected white components in its 1-neighborhood. In the former case the pixel forms the end of a path, otherwise it is the meeting point of at least three edges.

Note that due to this step, the maximum degree of the graphs we detect is limited to four. This is inevitable if nodes are detected at single-pixel locations. For higher degree nodes we will create several nodes of limited degree that are very close to each other and can be merged by a later post-processing step.

Given the node positions, it is very simple to find the edges. We perform a variant of breadth first search on the white pixels in the skeleton, starting from each node simultaneously. Each white pixel around a node gets a unique number and a queue. In each step we iterate over all queues and take out the first pixel. If it is unmarked, we mark it with the unique number of this queue and enqueue all its white neighbors. Otherwise, we have detected an edge, i.e. there is a path along white pixels that connects two nodes.

While walking along the pixels we record the length of the edge. Horizontal and vertical steps count as one unit, diagonal steps count as $\sqrt{2} \approx 1.41$ units.

The diameter of an edge calculated by computing a distance transform on the

segmented image. This assumes that the thinned edge lies in the middle of the actual edge. Computing the diameters is now a simple lookup of each edge-pixel from the skeleton in the distance transformed image. As we have the diameters along the whole edge on hand by this procedure, we then compute a median and a variance.

For handling the graph we rely on [NetworkX \[11\]](#).

1.3.4 Graph Filter Collection

The graph filter collection offers the possibility to add powerful processing steps that directly apply to the graph obtained after graph detection.

Often it is possible to improve the result by removing unwanted artifacts in the segmented image or during later processing stages. A common strategy, used for example in [\[4\]](#), [\[5\]](#), consists of “reparing” the errors in the skeleton using heuristics or user assisted methods. However, these methods carry the potential danger of introducing additional errors.

NEFI pursues a novel approach by exploiting the structure of the extracted graph. First, we retain a maximum of structural information by not altering the segmented image or the skeleton at all, i.e. we establish the graph including all artifacts. Then we use dedicated graph filters to remove said artifacts. For example, if the network in the input image is reasonably large, it will result in a large connected component in the graph. Small components resulting from noise can thus be removed effectively and safely. Since the effects of filtering the graph can immediately be evaluated by visual inspection, we prefer graph filtering over less transparent approaches that take place before the graph was established. [Figure 1.2](#) illustrates the use of filtering.

We have used filtering with sensitive segmentation to obtain surprisingly good results. Overly sensitive segmentation picks up fine detail but also introduces artifacts. However, almost all of the artifacts result in very small components that can easily be removed by filtering. The desired detail will remain mostly unaffected because it is part of the largest component. The graph depicted in [Figure 1.5](#) was obtained using this technique.

Filtering may also be used to remove parts of the graph which are not of interest. The following filters are predefined in NEFI. A filter removing everything not in the largest connected component, one smoothing vertices of degree two except if this introduces parallel edges and finally, one which removes all vertices and edges that are not contained in a cycle. Filters may be freely combined in any order. Naturally, the filter collection is designed for extendability.

Graph filtering and its various applications delivers excellent results. To our knowledge, no other software offers such tools as part of its core workflow. For this reason, we consider this one of NEFI’s strong points.

Evaluation

To assess NEFI's performance we investigate the output quality given input images of varying quality. Additionally we report on NEFI's speed.

1.3.5 Defining a Graph Similarity Measure

As a quality measure, we need to quantify the degree of congruence between the graph depicted in the original input image i and the graph computed by NEFI.

Let $A = (V_A, E_A)$ be the *true* graph correctly describing the structure depicted in i , with V_A and E_A denoting its vertex and edge set respectively. We call A the *ground truth*, which is of course not known in general. Furthermore, let B denote the graph obtained by executing one of NEFI's pipelines. Note that, $A, B \in \mathcal{G}$, where \mathcal{G} denotes the set of undirected edge-weighted planar graphs where vertices are labeled with their respective euclidean coordinates in the plane. With these definitions we propose a similarity measure s mapping any pair of graphs $A, B \in \mathcal{G}$ onto a number $s \in [0, 1]$.

We compute a correspondence of vertices in A to vertices in B . Two edges $e \in E_A$ and $f \in E_B$ then correspond if their endpoints correspond. We choose the correspondence such that the following intuitive notion of similarity are optimized.

1. Positions of vertices in V_A are similar to positions of corresponding vertices in V_B .
2. Edges in E_A including their weights are similar corresponding edges in E_B .

For an exact definition of s and the notions of *similarity* and *correspondence*, we refer the reader to the supplementary material.

We require that the measure is maximal if any graph A is compared with itself, that is $s(A, A) = 1$. Consequently, if A is completely different from B we have $s(A, B) = 0$. This minimum value is assumed if no viable correspondence between V_A and V_B can be found. Naturally, the value of $s(A, B)$ increases (decreases) if the similarity between A and B increases (decreases).

Transfer this
part to the
main text ...

--

--

--

Bibliography

- [1] “A threshold selection method from gray-level histograms”, *Systems, Man and Cybernetics, IEEE Transactions on*, vol. 9, no. 1, pp. 62–66, Jan. 1979, ISSN: 0018-9472. DOI: [10.1109/TSMC.1979.4310076](https://doi.org/10.1109/TSMC.1979.4310076).
- [2] M. Bastian, S. Heymann, and M. Jacomy, “Gephi: an open source software for exploring and manipulating networks”, 2009. [Online]. Available: <http://www.aaai.org/ocs/index.php/ICWSM/09/paper/view/154>.
- [3] V. Batagelj and A. Mrvar, “Pajek-program for large network analysis”, *Connections*, vol. 21, no. 2, pp. 47–57, 1998.
- [4] W. Baumgarten and M. J. Hauser, “Detection, extraction, and analysis of the vein network”, *Journal of Computational Interdisciplinary Sciences*, vol. 1, no. 3, pp. 241–249, 2010.
- [5] —, “Computational algorithms for extraction and analysis of two-dimensional transportation networks”, *J. Comput. Interdiscip. Sci*, vol. 3, pp. 107–16, 2012.
- [6] G. Bradski, “The OpenCV Library”, *Dr. Dobb’s Journal of Software Tools*, 2000.
- [7] D. Chai, W. Forstner, and F. Lafarge, “Recovering line-networks in images by junction-point processes”, in *Computer Vision and Pattern Recognition (CVPR), 2013 IEEE Conference on*, IEEE, 2013, pp. 1894–1901.
- [8] M. T. Dehkordi, S. Sadri, and A. Doosthoseini, “A review of coronary vessel segmentation algorithms”, *Journal of medical signals and sensors*, vol. 1, no. 1, p. 49, 2011.
- [9] Z. Guo and R. W. Hall, “Parallel thinning with two-subiteration algorithms”, *Communications of the ACM*, vol. 32, no. 3, pp. 359–373, 1989.
- [10] A. A. Hagberg, D. A. Schult, and P. J. Swart, “Exploring network structure, dynamics, and function using NetworkX”, in *Proceedings of the 7th Python in Science Conference (SciPy2008)*, Pasadena, CA USA, Aug. 2008, pp. 11–15.
- [11] —, “Exploring network structure, dynamics, and function using NetworkX”, in *Proceedings of the 7th Python in Science Conference (SciPy2008)*, Pasadena, CA USA, Aug. 2008, pp. 11–15.

- [12] M. Krause, R. M. Alles, B. Burgeth, and J. Weickert, “Fast retinal vessel analysis”, *Journal of Real-Time Image Processing*, pp. 1–10, 2013.
- [13] J. Leskovec and R. Sosič, *SNAP: a general purpose network analysis and graph mining library in C++*, <http://snap.stanford.edu/snap>, Jun. 2014.
- [14] S. Loscalzo and L. Yu, “Social network analysis: tasks and tools”, in *Social computing, behavioral modeling, and prediction*, Springer, 2008, pp. 151–159.
- [15] F. Meyer, “Un algorithme optimal pour la ligne de partage des eaux”, *Dans 8me congrès de reconnaissance des formes et intelligence artificielle*, vol. 2, pp. 847–857, 1991.
- [16] M. Negri, P. Gamba, G. Lisini, and F. Tupin, “Junction-aware extraction and regularization of urban road networks in high-resolution sar images”, *Geoscience and Remote Sensing, IEEE Transactions on*, vol. 44, no. 10, pp. 2962–2971, Oct. 2006, ISSN: 0196-2892. DOI: [10.1109/TGRS.2006.877289](https://doi.org/10.1109/TGRS.2006.877289).
- [17] M. Newman, “The structure and function of complex networks”, *SIAM Review*, vol. 45, no. 2, pp. 167–256, 2003. DOI: [10.1137/S003614450342480](https://doi.org/10.1137/S003614450342480). eprint: <http://dx.doi.org/10.1137/S003614450342480>. [Online]. Available: <http://dx.doi.org/10.1137/S003614450342480>.
- [18] B. Obara, M. Frickerc, and V. Graua, “Contrast independent detection of branching points in network-like structures”, in *SPIE Medical Imaging*, International Society for Optics and Photonics, 2012, pp. 83141L–83141L.
- [19] B. Obara, V. Grau, and M. D. Fricker, “A bioimage informatics approach to automatically extract complex fungal networks”, *Bioinformatics*, vol. 28, no. 18, pp. 2374–2381, 2012.
- [20] C. Rother, V. Kolmogorov, and A. Blake, “Grabcut - interactive foreground extraction using iterated graph cuts”, *ACM Trans. Graphics (SIGGRAPH)*, pp. 309–314, 2004.
- [21] E. Samuel, C. de la Higuera, and J.-C. Janodet, “Extracting plane graphs from images”, English, in *Structural, Syntactic, and Statistical Pattern Recognition*, ser. Lecture Notes in Computer Science, E. R. Hancock, R. C. Wilson, T. Windeatt, I. Ulusoy, and F. Escolano, Eds., vol. 6218, Springer Berlin Heidelberg, 2010, pp. 233–243, ISBN: 978-3-642-14979-5. DOI: [10.1007/978-3-642-14980-1_22](https://doi.org/10.1007/978-3-642-14980-1_22). [Online]. Available: http://dx.doi.org/10.1007/978-3-642-14980-1_22.
- [22] K. Xu, C. Tang, R. Tang, G. Ali, and J. Zhu, “A comparative study of six software packages for complex network research”, in *Communication Software and Networks, 2010. ICCSN '10. Second International Conference on*, Feb. 2010, pp. 350–354. DOI: [10.1109/ICCSN.2010.34](https://doi.org/10.1109/ICCSN.2010.34).

- [23] T. Zhang and C. Y. Suen, “A fast parallel algorithm for thinning digital patterns”, *Communications of the ACM*, vol. 27, no. 3, pp. 236–239, 1984.

--

--

--

Todo list

Transfer this part to the main text ... 19